



UNIVERSITY OF
GOTHENBURG



CHALMERS
UNIVERSITY OF TECHNOLOGY

MapReduce programming

DAT470/DIT065 Computational techniques for large-scale data

2024-04-15

Matti Karppa

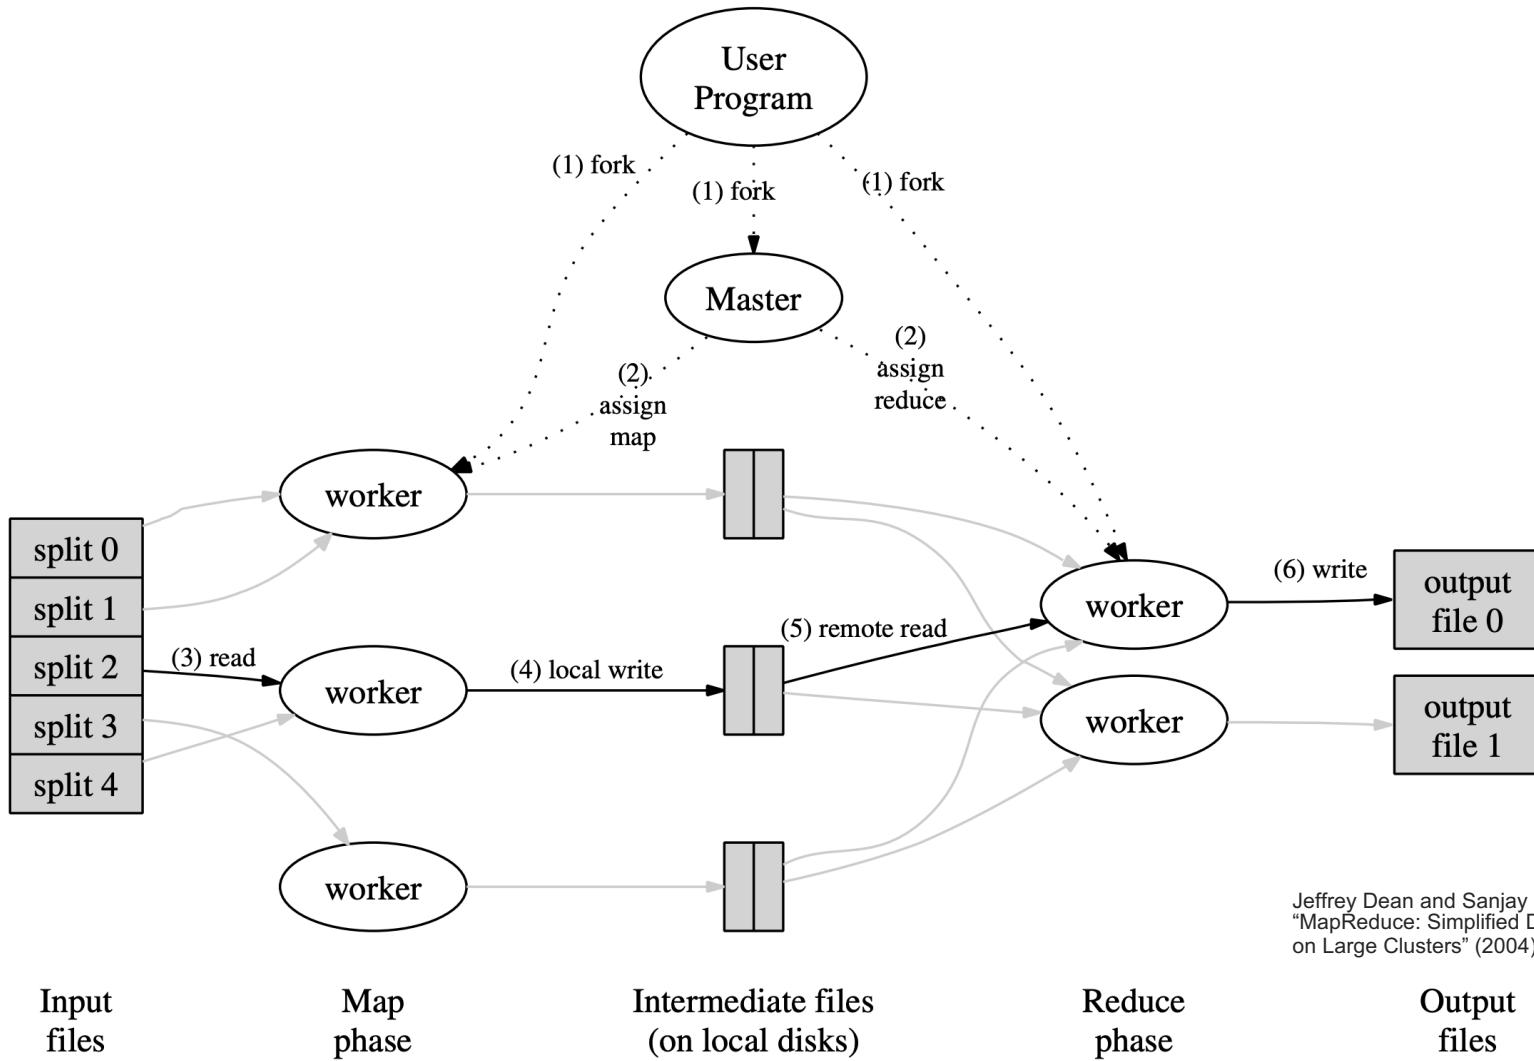


Motivation

- Computing on a large number of computers has challenges
 - Data has to be stored in a distributed fashion: no individual node can hold all data
 - Computing is **error-prone**: computers fail
 - **I/O bottleneck** presents a challenge: moving data from one compute node to another is very expensive
 - Coordinating computation is very difficult and laborious
- **MapReduce** is a solution that attempts to solve these some of these problems by assuming a (somewhat rigid) model of computation
 - Computation is **brought to data**: reduces the amount of I/O
 - Computation is modelled operations that adhere to a strict pattern, so they are easy to schedule
 - Likewise, if we detect a fault, we can simply replay those operations that failed

Map & Reduce

- In MapReduce programming, all data is treated as **key-value** pairs
- Computations are composed of two kinds of operations: **map** and **reduce** operations
- Map: $(K_1, V_1) \rightarrow \text{List}((K_2, V_2))$
 - Each key-value pair is mapped to a one or more key-value pairs
 - Maps are executed in parallel on different nodes (wherever the data lies)
 - Cf. Python's `map`
- After the map, the data is **shuffled** by hashing the keys
 - This (pseudorandomly) distributes the output to different nodes
 - Each key-value pair sharing the same key is guaranteed to end on the **same node**
- Reduce: $(K_2, \text{List}(V_2)) \rightarrow \text{List}((K_3, V_3))$
 - All values that were mapped to the same key are given as a list
 - **Typically**, reduce would then compute a single value from that, but can in practice return multiple key-value pairs
 - The function should be **commutative** and **associative** because the order of values cannot be guaranteed
 - Cf. Python's `reduce`





Example

- Let us count the frequency of each letter as in the previous lecture
- Data will be split by the lines; the input key is not significant (it can be a dummy key or line number or something of the sort)
- Map:
 - Key: *ignored*, value: line of text
 - For each character c in the line, emit $(c, 1)$
 - “We saw c once”
- Data will be shuffled to different nodes; all 1s of the same character will end up on the same node
- Reduce:
 - Key: c , value: list v_1, v_2, \dots of 1s, corresponding to the count
 - Emit $(c, \sum v_i)$

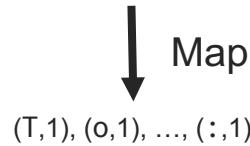
Execution example

- Suppose node 1 holds the line To be, or not to be, that is the question: and node 2 holds the line Whether 'tis nobler in the mind to suffer
- The **mapper** of node 1 emits (T,1), (o,1), ..., (:,1)
- The mapper of node 2 emits (W,1), (h,1), ..., (r,1)
- Data is then shuffled
 - Suppose e ends up on node 1 and t on node 2
 - Five pairs (e,1) are communicated from node 2 to node 1
 - Six pairs (t,1) are communicated from node 1 to node 2
- For the key e, the reducer on node 1 takes in (e,[1,1,...]), a list of 9 ones
- Reducer on node 1 emits (e,9)
- For the key t, the reducer on node 2 takes in (t,[1,1,...]), a list of 10 ones
- Reducer on node 2 emits (t,10)

Execution example

Node 1

To be, or not to be, that is the question:



$(t,1), (t,1), (t,1),$
 $(t,1), (t,1), (t,1) \text{ etc.}$

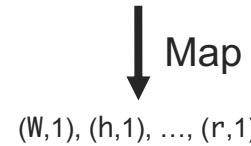
$(e,[1, 1, \dots, 1])$

Reduce

$(e,9)$

Node 2

Whether 'tis nobler in the mind to suffer



$(e,1), (e,1), (e,1),$
 $(e,1), (e,1) \text{ etc.}$

$(t,[1, 1, \dots, 1])$

Reduce

$(t,10)$



- Apache Hadoop is a free software ecosystem for big data processing
- Hadoop MapReduce is the standard implementation of MapReduce (originally introduced by Google)
- The ecosystem contains other supporting components, such as Hadoop Distributed Filesystem (HDFS), a fault-tolerant distributed filesystem that we'll discuss on the next lecture



MRJob

- MRJob is a Python library for implementing MapReduce jobs
- Works like this: the user implements a subclass of job class and implements the required operations
- Jobs can be run locally (with the “local runner”, this is the approach we’ll adopt in our assignments), or through the Hadoop infrastructure
- Specialty: The Python file implementing the class **must not contain** anything else
 - This is because the job is sent to Hadoop for running and this sets limitations on, e.g., how the Python code may interact with its environment

Generator functions

- The functions in MRJob are written as **generator functions**
- These functions are syntactic sugar for implementing **lazy functions** that behave as if **iterators**
- In Python, an object is **iterable** if it implements `__iter__` and `__next__`, that is, it allows one to create an **iterator** and (typically) go through the data structure one element at a time by advancing the iterator
 - All basic data structures in the Python standard library are iterable
- The functions `__iter__` and `__next__` are usually called via the convenience functions `iter` and `next`, respectively
 - Suppose X is some kind of collection
 - `it = iter(X)` returns an iterator to the collection
 - `next(it)` returns the next element in the collection and advances the iterator; if all elements have been returned, it raises the `StopIteration` exception instead
- The Python `for x in X` loop is essentially syntactic sugar for the following construct:
 - Create an iterator with `it = iter(X)`
 - Get the next element with `x = next(it)`
 - Catch the `StopIteration` exception and break from loop

Example

- These two are functionally the same

```
for x in X:  
    print(x)
```

```
it = iter(X)  
while True:  
    try:  
        x = next(it)  
        print(x)  
    except StopIteration:  
        break
```

Generator functions

- If a Python function contains a `yield` keyword, it is implicitly wrapped into an iterator object!
 - The next value returned by `__next__` is determined by `yield`
 - When the function ends executing, it raises `StopIteration`
 - This means that we can say `x in f()`
- Advantages:
 - From the caller's point of view, the function behaves like a collection even though the values do not exist at the moment of call
 - This enables us to lazily construct the values that the caller wants only upon need
 - This can save **a lot of** space and time if we do not need access to all values at one time, or if we abort early
- Disadvantages:
 - This introduces tacitly new stateful objects which we need to be aware of
 - This can be inefficient if used inconsiderately

Example: return the n first Fibonacci numbers

```
def fibonacci(n):
    x = 0
    y = 0
    z = 1
    for _ in range(n):
        x = y
        y = z
        z = x+y
        yield y

if __name__ == '__main__':
    for x in fibonacci(10):
        print(x)
```

Example: Letter count with MRJob

```
from mrjob.job import MRJob
class MRLettercount(MRJob):
    def mapper(self, _, line):
        for c in line:
            yield (c,1)
def reducer(self, c, ones):
    yield (c,sum(ones))
if __name__ == '__main__':
    MRLettercount.run()
```

We ignore the key

This class defines a MapReduce job; we extend it

The map function is defined by overriding the mapper function

Values are returned lazily; the function is actually a generator

The reducer works just the same

The default runner reads data from a file or stdin and provides standard command-line arguments; there must be **nothing** else in the .py file

Example: Letter count with MRJob

```
(mrjob) karppa@CM-DWQ77WNX67 lecture6 % cat lines.txt | tee /dev/tty | python3 mrjob_lettercount1.py
```

```
-r local --num-cores 4 2> /dev/null
```

```
To be, or not to be, that is the question:
```

```
Whether 'tis nobler in the mind to suffer
```

| | |
|-----|----|
| "q" | 1 |
| "r" | 4 |
| "s" | 4 |
| "t" | 10 |
| "e" | 9 |
| "f" | 2 |
| "h" | 5 |
| "i" | 5 |
| "l" | 1 |
| "m" | 1 |
| "n" | 5 |
| "o" | 7 |
| "'" | 1 |
| "," | 2 |
| 1 | |
| "T" | 1 |
| "W" | 1 |
| "a" | 1 |
| "b" | 3 |
| "d" | 1 |
| "u" | 2 |
| " " | 16 |

Protip: print the intermediate status on screen by writing to the terminal file

Use “local” runner (default is debug runner which does not scale)

Set the number of cores (workers)

Protip: MRJob writes a lot of stuff into stderr; this suppresses it for clearer output for demonstration purposes

Looks about right



Performance considerations

- MapReduce only makes sense if the data is sufficiently large such that it cannot neatly fit in RAM
 - Then the framework solves the problem of coordinating work (although in a somewhat rigid way)
 - Computation is brought to the data, as compute is faster than I/O
 - If the data fits in RAM, then MapReduce only introduces unnecessary overhead
- The key to a performant MapReduce algorithm is to limit the amount of I/O
 - In practice: limit the amount of data that is communicated in the shuffling stage
- Storage of data matters
 - If the data is stored evenly distributed, we're good, but skewed distribution (e.g., all relevant pieces on one node) hurts parallelization



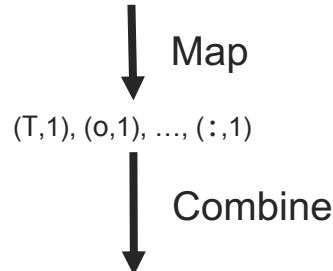
Combiners

- A problem with our example is that it can generate **a lot of** communication
 - Each occurrence of a letter is potentially communicated separately to another node
 - However, we only care about the **sum**: would it not be more efficient to just communicate the sum of counts on a node?
- **Combiners** solve this problem
- Combiners are essentially “mini-reducers” that are executed **locally** after the map operation but before shuffling
- Combiners perform a reduce operation locally, and only the result is then communicated
 - E.g., in this case, the combiner could simply count the sum of the occurrences locally

Execution example

Node 1

To be, or not to be, that is the question:



($e,4$) ($t,6$)

($t,6$)

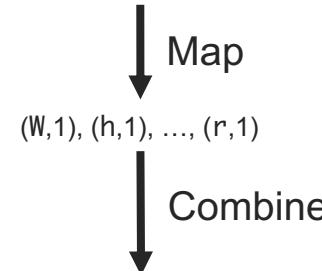
($e,[4,5]$)

Reduce

($e,9$)

Node 2

Whether 'tis nobler in the mind to suffer



($e,5$) ($t,4$)

($e,5$)

($t,[6,4]$)

Reduce

($t,10$)



Example

```
from mrjob.job import MRJob

class MRLettercount(MRJob):
    def mapper(self, _, line):
        for c in line:
            yield (c,1)

    def combiner(self, c, ones):
        yield (c,sum(ones))

    def reducer(self, c, counts):
        yield (c,sum(counts))

if __name__ == '__main__':
    MRLettercount.run()
```

```
(mrjob) karppa@CM-DWQ77WNX67 lecture6 % cat lines.txt | tee /dev/tty | python3 mrjob_lettercount2.py  
-r local --num-cores 4 2> /dev/null
```

To be, or not to be, that is the question:

Whether 'tis nobler in the mind to suffer

| | |
|-------|----|
| "o" | 7 |
| "q" | 1 |
| "r" | 4 |
| "b" | 3 |
| "d" | 1 |
| "e" | 9 |
| "f" | 2 |
| "h" | 5 |
| "i" | 5 |
| "l" | 1 |
| "m" | 1 |
| "n" | 5 |
| 1 | |
| "T" | 1 |
| "W" | 1 |
| "a" | 1 |
| "s" | 4 |
| "t" | 10 |
| "u" | 2 |
| " " | 16 |
| "'" | 1 |
| " , " | 2 |



Multiple steps

- What if we wanted to figure out the minimum, maximum, and average counts of characters?
- What if we wanted to figure out **which** characters have the minima and maxima counts?
- We can solve this by creating **multiple steps**
 - Each step is a MapReduce operation that is applied to the key-value pairs output by the previous step
- We change the first step such that we output the count of characters as **values**:
(character,count)-tuples; we use a dummy key
- We don't really need a mapper in the second step as we only do another reduce step
(mapper can be thought to be as the identity):
 - We traverse through the lists and identify the minima and maxima
- We return multiple key-value pairs to identify the identities and the values requested

Execution example



UNIVERSITY OF
GOTHENBURG



CHALMERS
UNIVERSITY OF TECHNOLOGY

Node 1

To be, or not to be, that is the question:

Map

(T,1), (o,1), ..., (:,1)

Combine

(e,4) (t,6)

Shuffle

(e,[4,5])

Reduce

(None,(e,9))

Map

(None,(e,9))

(None,(e,9))

Reduce

(min,(a,1)), (max,(t,10))

Node 2

Whether 'tis nobler in the mind to suffer

Map

(W,1), (h,1), ..., (r,1)

Combine

(e,5) (t,4)

(e,5)

(t,[6,4])

Reduce

(None,(t,10))

Map

(None,(t,10))

(None,(t,10))

Step 1

Step 2



Steps in MRJob

- Multiple steps can be created by
 1. Implement all operations required as distinct member functions
 2. Implement the `steps()` function that returns a list of `MRStep` objects (each object defines the mapper, combiner, and reducer for the step)
 3. That's it!



```
from mrjob.job import MRJob
from mrjob.step import MRStep

class MRLettercount(MRJob):
    def mapper(self, _, line):
        for c in line:
            yield (c,1)

    def combiner(self, c, ones):
        yield (c,sum(ones))

    def reducer(self, c, counts):
        yield (None, (c,sum(counts)))

    def reducer2(self, _, charcounts):
        minimum = (None, float('inf'))
        maximum = (None, -1)
        numchars = 0
        total = 0
        for (c,count) in charcounts:
            numchars += 1
            total += count
            if count < minimum[1]:
                minimum = (c,count)
            if count > maximum[1]:
                maximum = (c,count)
        yield ('min', minimum)
        yield ('max', maximum)
        yield ('avg', total / numchars)

    def steps(self):
        return [MRStep(mapper=self.mapper, combiner=self.combiner,
                      reducer=self.reducer),
                MRStep(reducer=self.reducer2)]
```



Example

```
(mrjob) karppa@CM-DWQ77WNX67:~/lecture6 % cat lines.txt | tee /dev/tty | python3 mrjob_lettercount3.py  
-r local --num-cores 4 2>/dev/null  
To be, or not to be, that is the question:  
Whether 'tis nobler in the mind to suffer  
"min"    ["q",1]  
"max"    [",",16]  
"avg"    3.772727272727273
```

Command line parameters

- MRJob provides a basic `argparse`-compatible command-line interface
- Suppose we want to make it optional to only include ASCII characters in our word count, how would we do this?
- We implement the member function `configure_args()` and add a `passthrough` argument
- The argument is then available through `self.options`

Example

```
class MRLettercount(MRJob):  
    def configure_args(self):  
        super(MRLettercount, self).configure_args()  
        self.add_passthru_arg(  
            '--ascii-only', default=False, action='store_true',  
            help='Only apply to ASCII characters')  
  
    def mapper(self, _, line):  
        if self.options.ascii_only:  
            for c in line:  
                if ord(c) < 128:  
                    yield (c,1)  
        else:  
            for c in line:  
                yield (c,1)
```

Required so that ordinary options are added as usual (call superclass version)

This adds the actual argument; note that this is exactly the same as using argparse

Args dictionary is available through self.options; it is exactly like the output of parse_args



Running MRJob programmatically

- It was already mentioned that MRJob class files must not contain any other code
- This is related to the way they are handled by Hadoop
- If you want to run MRJob jobs manually (e.g., to instrumentate the code to be able to measure runtimes!), you will need to create a custom runner
- Do as follows:
 1. Construct your object, use the constructor argument `args` to pass the command-line parameters as usual (this can include your custom parameters)
 2. Construct the runner object with the member function `make_runner()` (idiom: use the `with` construction to gracefully close the object when it's finished)
 3. Use the `run()` member function of the runner to actually run the task
 4. Recover output using the `parse_output()` and `cat_output()` functions

Example

Construct the job

```
from mrjob_lettercount4 import MRLettercount
import time
import argparse

if __name__ == '__main__':
    parser = argparse.ArgumentParser()
    parser.add_argument('-w', '--num-workers', type=int, default=1)
    parser.add_argument('filename')
    args = parser.parse_args()
    mr_job = MRLettercount(args=[ '-r', 'local', '--num-cores',
        str(args.num_workers), '--ascii-only',
        args.filename])
```

Use local runner (do this on bayes as well)

Set the number of cores

```
Create runner, it will be closed gracefully when it goes out of scope from the with block
start = time.time()
with mr_job.make_runner() as runner: Run the job
    runner.run()
    for key, value in mr_job.parse_output(runner.cat_output()):
        print(key,value)
end = time.time()
print(f'Number of workers: {args.num_workers}')
print(f'Time elapsed: {end-start} s')
```

Custom arguments

Raw text output

Parse output into key-value pairs



Example

```
(mrjob) karppa@CM-DWQ77WNX67 lecture6 % python3 mrjob_lettercount_runner.py -w 4 .../lecture4/shakespeare.txt
No configs specified for local runner
min ['\\', 1]
max [' ', 949869]
avg 61111.96511627907
Number of workers: 4
Time elapsed: 4.8429319858551025 s
```